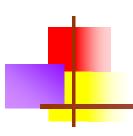


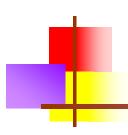
State-Space Searches





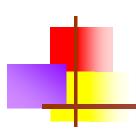
State spaces

- A state space consists of
 - A (possibly infinite) set of states
 - The start state represents the initial problem
 - Each state represents some configuration reachable from the start state
 - Some states may be goal states (solutions)
 - A set of operators
 - Applying an operator to a state transforms it to another state in the state space
 - Not all operators are applicable to all states
- State spaces are used extensively in Artificial Intelligence (AI)

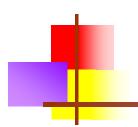


Example 1: Maze

- A maze can be represented as a state space
 - Each state represents "where you are" in the maze
 - The start state represents your starting position
 - The goal state represents the exit from the maze
- Operators (for a rectangular maze) are: move north, move south, move east, and move west
 - Each operator takes you to a new state (maze location)
 - Operators may not always apply, because of walls in the maze



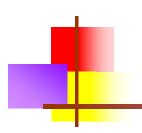
"You are given two jugs, a 4-litre one and a 3-litre one. Neither has any measuring markers on it. There is a pump that can be used to fill the jugs with water. How can you get exactly 2 litres of water into 4-litre jug."



 \blacksquare State: (x, y)

$$x = 0, 1, 2, 3, \text{ or } 4$$
 $y = 0, 1, 2, 3$

- \blacksquare Start state: (0, 0).
- Goal state: (2, n) for any n.
- Attempting to end up in a goal state.



1.
$$(x, y)$$

if $x < 4$

$$\rightarrow$$
 (4, y)

2.
$$(x, y)$$
 if $y < 3$

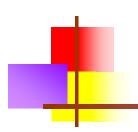
$$\rightarrow$$
 (x, 3)

3.
$$(x, y)$$
 if $x > 0$

$$\rightarrow$$
 (x - d, y)

4.
$$(x, y)$$
 if $y > 0$

$$\rightarrow$$
 (x, y - d)



5.
$$(x, y) \rightarrow (0, y)$$

if $x > 0$

6.
$$(x, y) \rightarrow (x, 0)$$

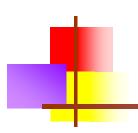
if $y > 0$

7.
$$(x, y) \rightarrow (4, y - (4 - x))$$

if $x + y \ge 4, y > 0$

8.
$$(x, y) \rightarrow (x - (3 - y), 3)$$

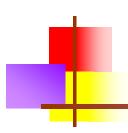
if $x + y \ge 3, x > 0$



9.
$$(x, y)$$
 $\rightarrow (x + y, 0)$
if $x + y \le 4$, $y > 0$
10. (x, y) $\rightarrow (0, x + y)$
if $x + y \le 3$, $x > 0$
11. $(0, 2)$ $\rightarrow (2, 0)$
12. $(2, y)$ $\rightarrow (0, y)$

State Space Search: Water Jug Problem

- 1. current state = (0, 0)
- 2. Loop until reaching the goal state (2, 0)
 - Apply a rule whose left side matches the current state
 - Set the new current state to be the resulting state
 - (0, 0)
 - (0, 3)
 - (3, 0)
 - (3, 3)
 - (4, 2)
 - (0, 2)
 - (2, 0)



Example 2: The 15-puzzle

Start state:

3	10	13	7
9	14	6	1
4		15	2
11	8	5	12

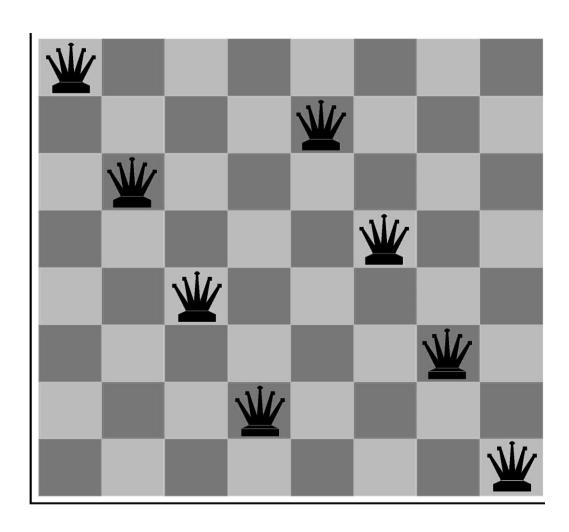
Goal state:

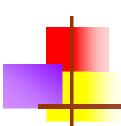
1	2	3	4
5	6	7	8
9	10	11	12
13	14	15	

- The start state is some (almost) random configuration of the tiles
- The goal state is as shown
- Operators are
 - Move empty space up
 - Move empty space down
 - Move empty space right
 - Move empty space left
- Operators apply if not against edge



Example 2: 8 Queen Problem

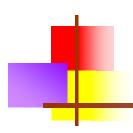




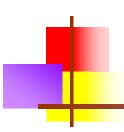
State-Space problem formulation

- states? -any arrangement of n<=8 queens
 -or arrangements of n<=8 queens in leftmost n columns, 1 per column, such that no queen attacks any other.

- <u>initial state?</u> no queens on the board
- actions? -add queen to any empty square
 -or add queen to leftmost empty square such that it is not attacked by other queens.
- goal test? 8 queens on the board, none attacked.
- <u>path cost?</u> 1 per move

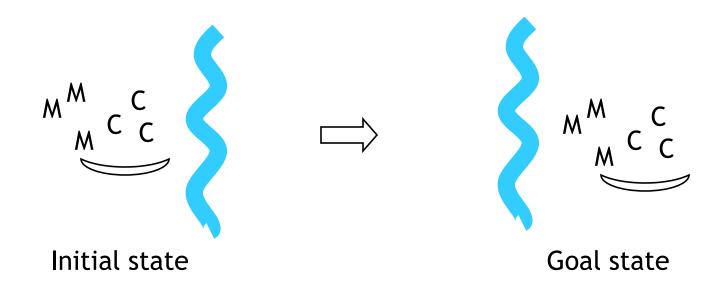


- Formuation 1
 - Operators: Add a queen in any square
- Formuation 2
 - Place the queen in the leftmost empty column
- Formulation 3



Example 3: Missionaries and cannibals

- An old puzzle is the "Missionaries and cannibals" problem (in various guises)
- The missionaries and cannibals wish to cross a river
- They have a canoe that can hold two people
- It is unsafe to have cannibals outnumber missionaries

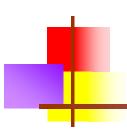


States

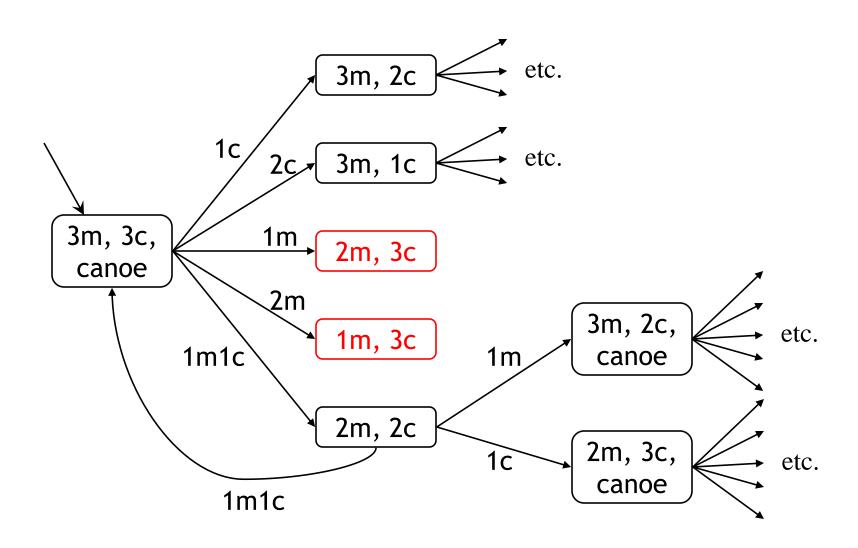
- A state can be represented by the number of missionaries and cannibals on each side of the river
 - Initial state: 3m,3c,canoe / 0m,0c
 - Goal state: 0m,0c / 3m,3c,canoe
 - We assume that crossing the river is a simple procedure that always works (so we don't have to represent the canoe being in the middle of the river)
- However, this is redundant; we only need to represent how many missionaries/cannibals are on one side of the river
 - Initial state: 3m,3c,canoe
 - Goal state: 0m,0c

Operations

- An *operation* takes us from one state to another
- Here are five possible operations:
 - Canoe takes 1 missionary across river (1m)
 - Canoe takes 1 cannibal across river (1c)
 - Canoe takes 2 missionaries across river (2m)
 - Canoe takes 2 cannibals across river (2c)
 - Canoe takes 1 missionary and 1 cannibal across river (1m1c)
- We don't have to specify "west to east" or "east to west" because only one of these will be possible at any given time

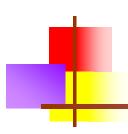


The state space



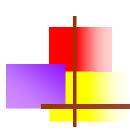
Example 3, revisited

- 3 missionaries, 3 cannibals, 1 canoe
- The canoe can hold at most two people
- Cannibals may never outnumber missionaries (on either side)
- Initial state is (3, 3, 1), representing the number of missionaries, cannibals, boats on the *initial* side
- The goal state is (0, 0, 0)
- Operators are addition or subtraction of the vectors
 (1 0 1), (2 0 1), (0 1 1), (0 2 1), (1 1 1)
- Operators apply if result is between (0 0 0) and (3 3 1)



State-space searching

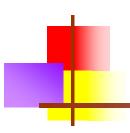
- Most problems in AI can be cast as searches on a state space
- The space can be tree-shaped or graph-shaped
 - If a graph, need some way to keep track of where you have been, so as to avoid loops
- The state space is often very, very large
- We can minimize the size of the search space by careful choice of operators
- Exhaustive searches don't work—we need *heuristics*



The basic search algorithm

Initialize: put the start node into OPEN while OPEN is not empty take a node N from OPEN if N is a goal node, report success put the children of N onto OPEN Report failure

- If OPEN is a stack, this is a depth-first search
- If OPEN is a queue, this is a breadth-first search
- If OPEN is a *priority queue*, sorted according to *most promising first*, we have a best-first search

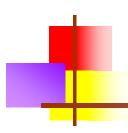


Heuristic searching

- All the previous searches have been blind searches
 - They make no use of any knowledge of the problem
 - If we know something about the problem, we can usually do much, much better
- Example: 15-puzzle
 - For each piece, figure out how many moves away it is from its goal position, if no other piece were in the way
 - The total of these gives a measure of distance from goal
 - This is a heuristic measure

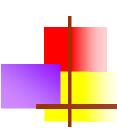


- A heuristic is a rule of thumb for deciding which choice might be best
- There is no general theory for finding heuristics, because every problem is different
- Choice of heuristics depends on knowledge of the problem space



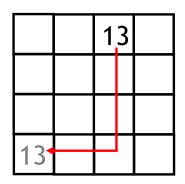
Best-first searching

- Use the same basic search algorithm
- Choose from OPEN the "best" node, that is, the one that seems to be closest to a goal
- Generally, even very poor heuristics are significantly better than blind search, but...
- No guarantee that the best path will be found
- No guarantee on the space requirements

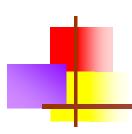


The 15-puzzle again

- Consider one piece in the 15-puzzle
 - If nothing were in the way, how many moves would it take to get this piece to where it belongs?



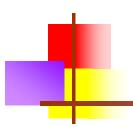
- This piece will have to be moved at least that many times to get it to where it belongs
- If we do this for every piece, and add up the moves, we get a (very) optimistic measure of how many moves it will take to solve the puzzle
- Suppose, from a given position, we try every possible single move (there can be up to four of them), and pick the move with the smallest sum
- This is a reasonable *heuristic* for solving the 15-puzzle



Iterative deepening

Set LIMIT to zero
do forever
Do a depth-first search up to LIMIT levels deep
If a goal is found, return success,
else add 1 to LIMIT

- Each time through the loop we start all over!
- If we find a path, it will be a shortest possible path
- Only requires linear space, because it only uses DFS
- Increased time requirements are only linear



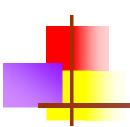
The A* algorithm

Suppose:

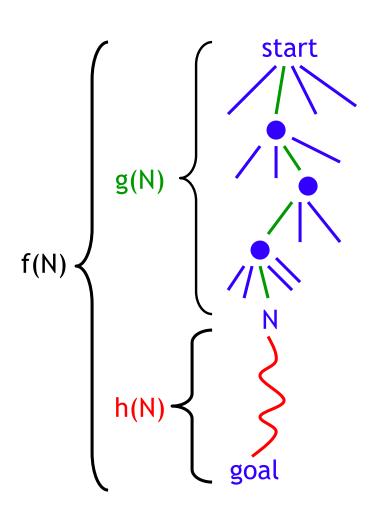
- You keep track of the distance g(N) that each node N that you visit is from the start state
- You have some heuristic function, h(N), that estimates the distance between node N and a goal

Then:

- f(N) = g(N) + h(N) gives you the (partially estimated) distance from the start node to a goal node
- The A* algorithm is: choose from OPEN the node N
 with the smallest value of f(N)



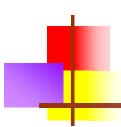
The A^* formula: f(N) = g(N) + h(N)



- g(N) is the (known)
 distance from start to N
- h(N) is the (estimated)
 distance from N to a goal
- f(N) is just the sum of these
- f(N) is our best guess as to the distance from start to a goal, passing through N

How good is A*?

- Memory usage depends on the heuristic function
 - If h(N) = constant, then $A^* = breadth-first$
 - If h(N) is a perfect estimator, A* goes straight to goal
 - ...but if h(N) were perfect, why search?
- Quality of solution also depends on h(N)
- It can be proved that, if h(N) is optimistic (never overestimates the distance to a goal), then A* will find a best solution (shortest path to a goal)



A* applied to the 15-puzzle

- Remember, if h(N) is optimistic (never overestimates the distance to a goal), then A* will find a best solution (shortest path to a goal)
- Our heuristic for the 15-puzzle was optimistic
- Therefore, using A* will find a solution with the fewest possible moves



- A* may require exponential storage
- Solution: Iterative-deepening A* (IDA*)
 - Just like Iterative Deepening, but...
 - ...instead of using g(N) (the actual depth so far) to cut off searching...
 - ...use f(N) (the estimated *total* depth)
- IDA* gives the same results as A*
- Since IDA* is basically a depth-first search, storage requirements are linear in length of path



- Many (or most) problems in AI can be represented as state-space searches
- The best searches combine a basic blind search technique with heuristic knowledge about the problem space
- A* and its variations, especially IDA*, are the best heuristic search techniques known

